

Scaling graph algorithm

(All Pair Shortest Path)

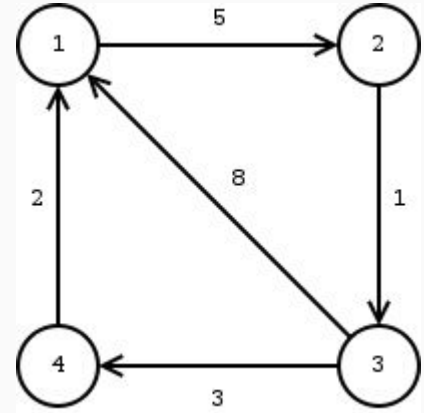
Project Members:-

Jatin Dev-18111027

Swapnil Raykar-18111078

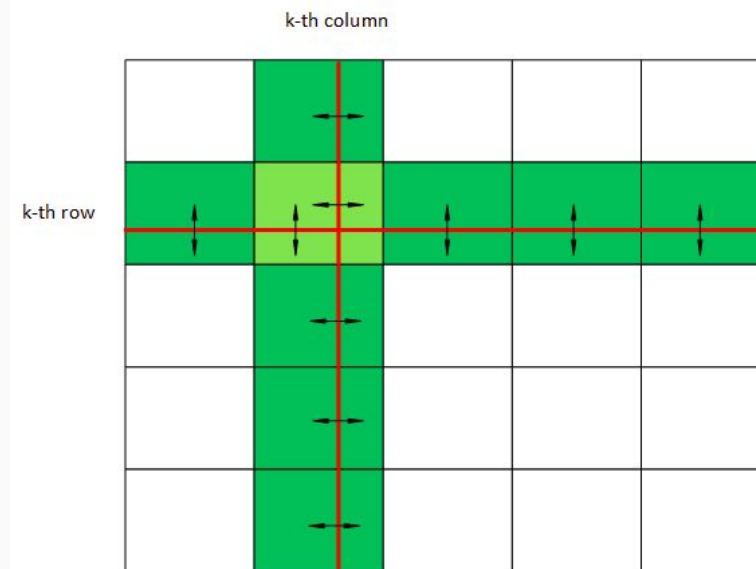
Introduction

- All Pair shortest path find shortest path between every pair of vertices in a graph.
- Applications:
 - Shortest paths in directed graphs
 - Optimal routing problem
 - Transitive Closure
 - Kleene's algorithm
- Parallelizing APSP problem can save lots of time.



Related Work

- **Speeding up APSP**
 - **Checkerboard:** The cost matrix P is divided into equal parts of size $(n/p) \times (n/p)$, and each is allocated to a different processor.



Related Work

- **Speeding up APSP**

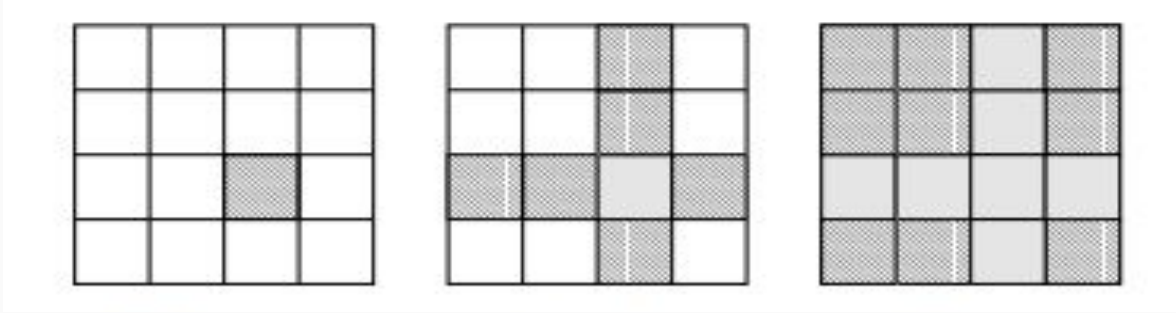
- **Checkerboard:** The cost matrix P is divided into equal parts of size $(n/p) \times (n/p)$, and each is allocated to a different processor.
- **Pipelined Checkerboard:** Remove the synchronisation step and each process starts the computation as soon as it has the data necessary to compute.
- **Blocked version APSP:** makes better utilization of the cache.
- Slight optimization on original APSP problem which gives $O(n^{2.4})$, $O((n^3)/\log n)$ algorithm.

Implementation

- We studied and implemented following algorithms:
 - Naive Checkerboard Version
 - Naive Parallel APSP
 - Blocked Parallel Version
 - Pipelined Blocked Parallel version

Implementation

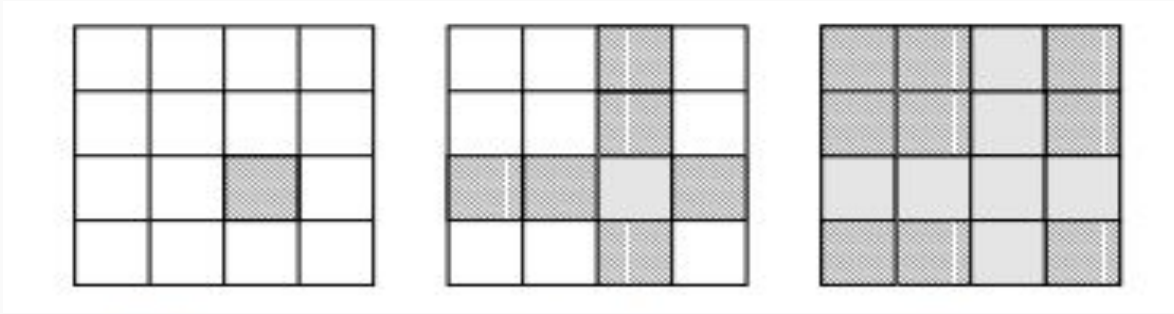
- **Blocked Parallel Version:**



First pivot element(2,2) compute its own block then it send it to corresponding column(2) and row(2) blocks. After their computation rest of the blocks are computed

Implementation

- **Blocked Parallel Version:**



First pivot element(**2,2**) compute its own block then it send it to corresponding column and row blocks. After their computation rest of the blocks are computed

- **Pipelined Blocked Version:** Removed barrier. Pivot block for the next iteration can't be blocked due to other processes.

Implementation

- **Input Graphs**

- **Random Graph Generator**

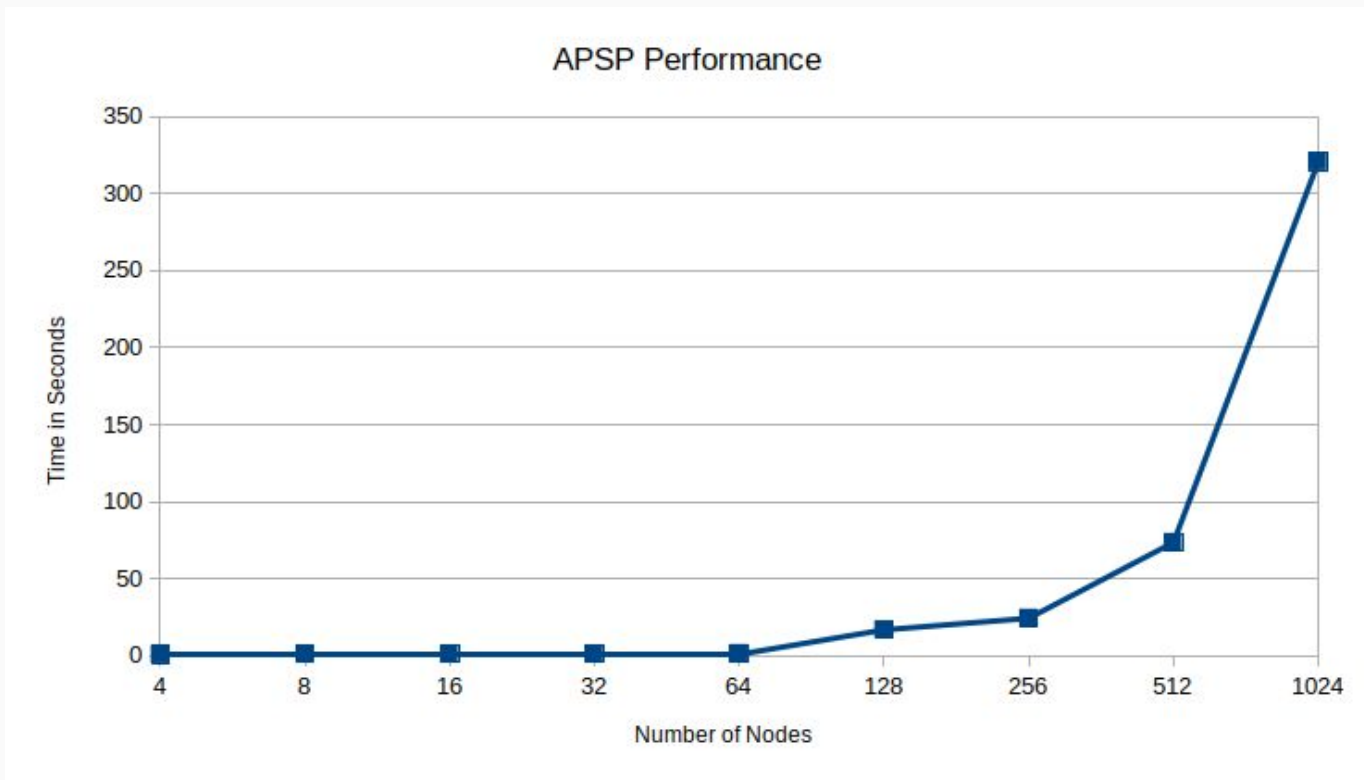
- python program

- **SNAP(Stanford Network Analysis Project) Graphs**

- C program to generate the graph matrix from the benchmark

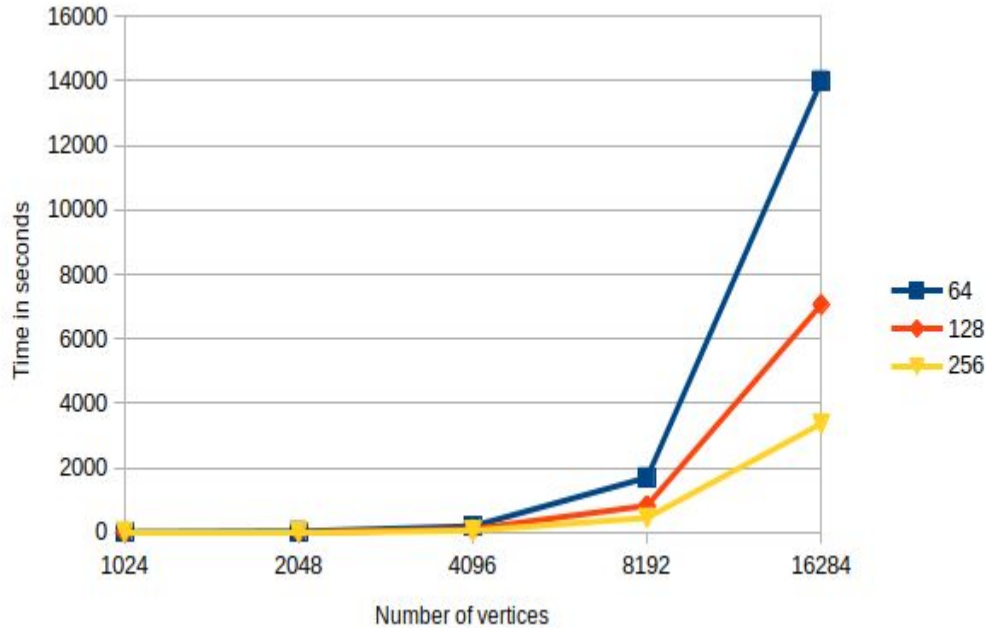
Experiments

Naive APSP



Experiments

Blocked Parallel Version - CSE Cluster

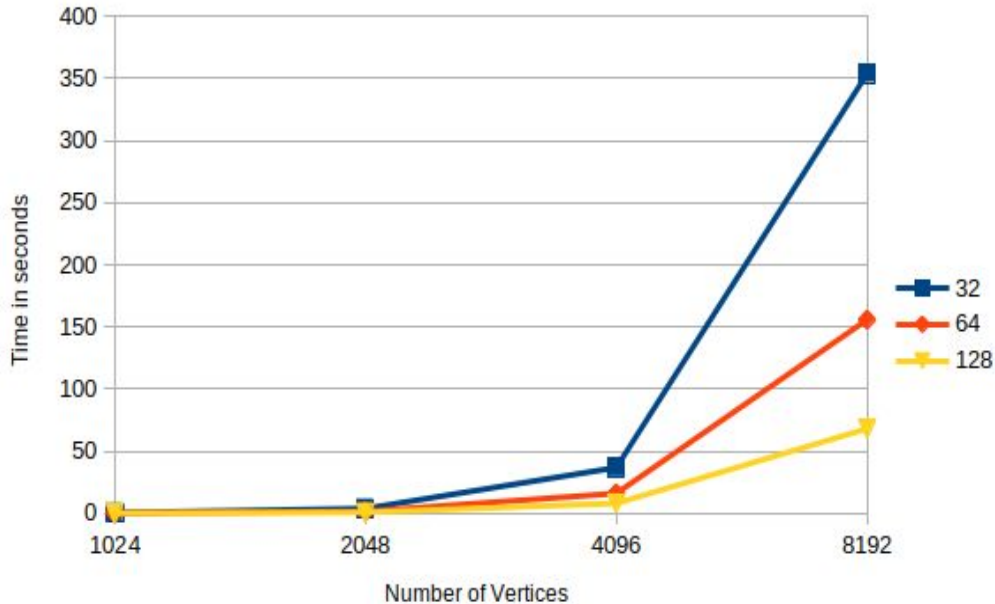


Observations:

- Time decreases ~50% with increase number of processes
- Decrease becomes more significant as graph size increases

Experiments

Blocked Parallel Version - HPC 2010

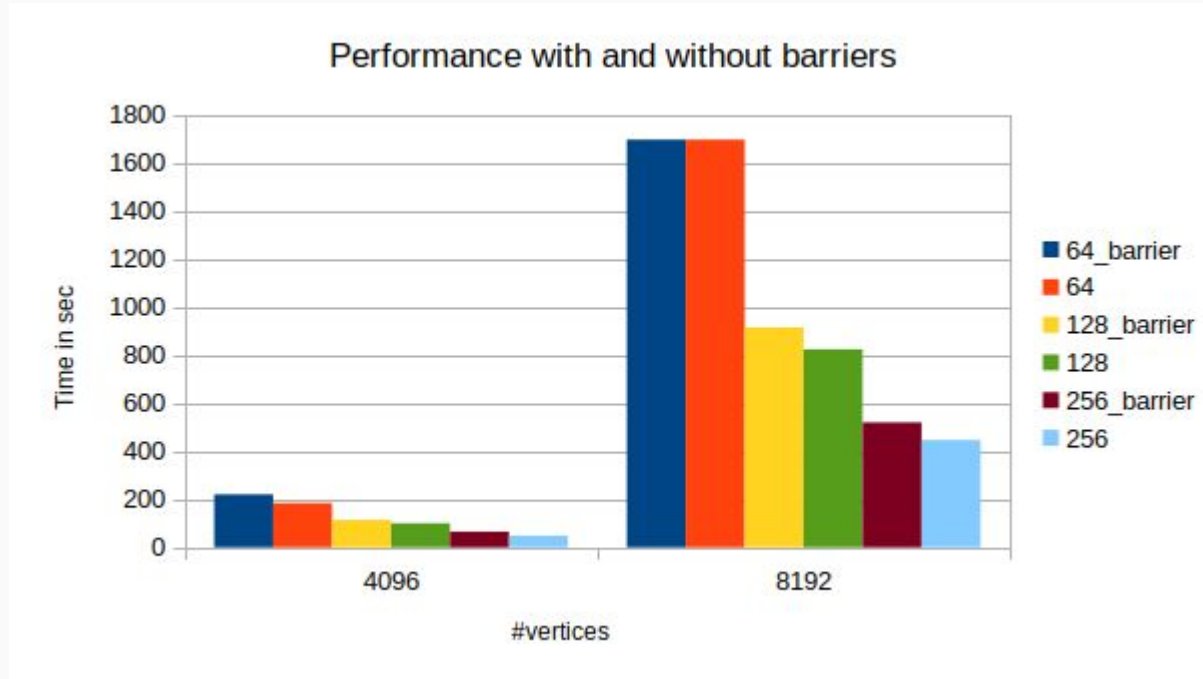


Observations:

- Time decreases with increase number of processes, becomes more significant with big graphs
- Less time required compared to CSE cluster

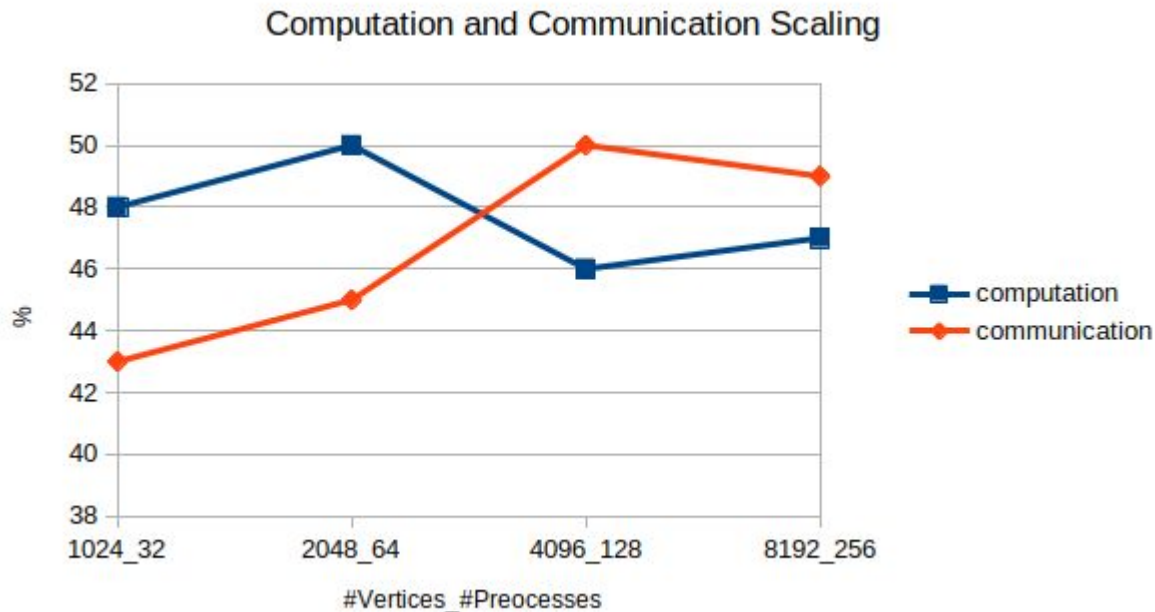
Performance Analysis

Pipelined Blocked Parallel Version



Performance Analysis

- TAU

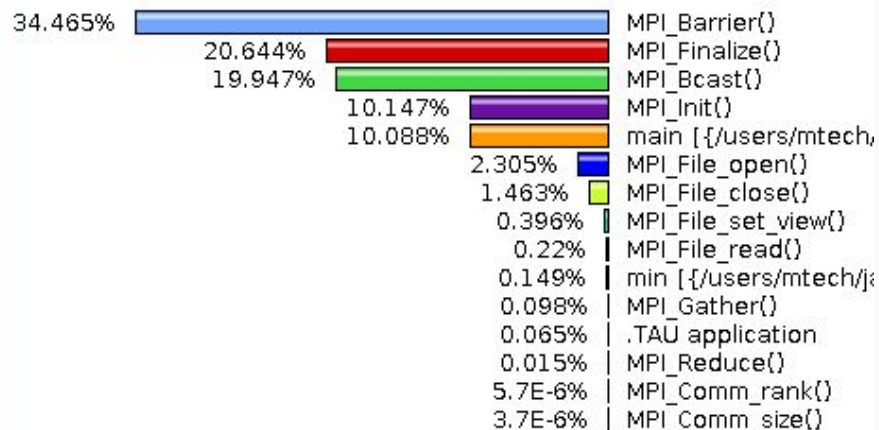


Performance Analysis

Blocked Parallel Version

Metric: TIME

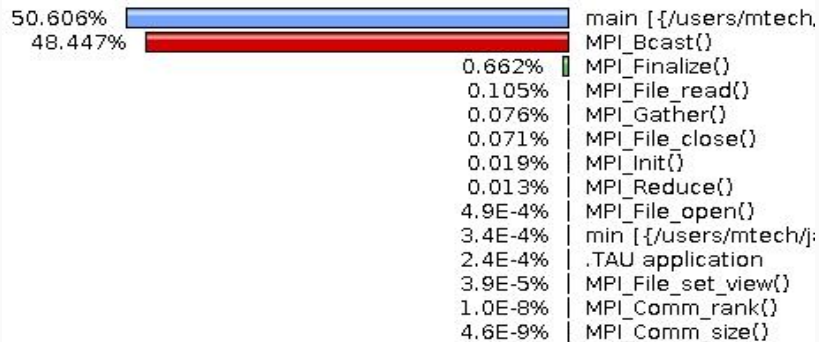
Value: Exclusive percent



Pipelined Blocked Parallel Version

Metric: TIME

Value: Exclusive percent



Profiling of 16384 vertices with 128 processes

Future work

- Assignment and placement of nodes
- Distribute load evenly amongst the nodes during a given iteration ,for example pivot node remains idle for most part of the iteration.
- Overlapping computation and communication,such that data of the next iteration becomes ready at the start of the previous iteration.

THANK YOU